

run_sjsdm_selected_candidate_folds.R

September 15, 2026

Contents

run_sjsdm_selected_candidate_folds	1
Index	3

run_sjsdm_selected_candidate_folds
<i>Run Selected sjSDM Candidate Across Folds</i>

Description

Refits one selected regularization candidate independently in every fold and materializes aligned out-of-fold probabilities.

Usage

```
run_sjsdm_selected_candidate_folds(  
  data_assignments = NULL,  
  data_selected_candidate = NULL,  
  data_sample_ids = NULL,  
  taxon_names = NULL,  
  prepare_fold_function = NULL,  
  fit_function = NULL,  
  predict_function = NULL,  
  seed = 900723L  
)
```

Arguments

- data_assignments**
Location-level assignment table accepted by [run_sjsdm_tuning_candidates()].
- data_selected_candidate**
One-row selected candidate table containing the candidate identifier, the six sjSDM regularization parameters, and regularization source.
- data_sample_ids**
Sample metadata with 'dataset_name' and 'age'. Optional sample, row, and location identifiers are derived when absent.

<code>taxon_names</code>	Unique character vector defining the full response-taxon order.
<code>prepare_fold_function</code>	Injectable function called with training and test row indices plus repeat and fold identifiers. It must return train/test inputs, retained train/test response matrices, the full test response matrix, aligned test sample identifiers, and the fold taxon mapping.
<code>fit_function</code>	Injectable function called with the prepared training input, selected candidate, and deterministic fit seed.
<code>predict_function</code>	Injectable function called with the fitted object and prepared test input. It must return probabilities aligned to retained test observations.
<code>seed</code>	Single non-negative integer used to derive stable fold fit seeds.

Details

The prediction table deliberately omits the candidate identifier. Selected candidate provenance remains in the separate fold diagnostics and selected candidate artifact. Taxa not retained in a training fold remain as rows with missing probabilities and their taxon-mapping status.

Value

Named list with data-prediction and data-diagnostic tibbles. Predictions contain one row per repeat, sample, and full-response taxon, including observed values, predictions, training-fold null probabilities, and explicit status. Diagnostics record one structured fit status per fold.

Examples

```
## Not run:
run_sjsdm_selected_candidate_folds(
  data_assignments = data_assignments,
  data_selected_candidate = data_selected_candidate,
  data_sample_ids = data_sample_ids,
  taxon_names = taxon_names,
  prepare_fold_function = prepare_fold_function,
  fit_function = fit_function,
  predict_function = predict_function
)

## End(Not run)
```

Index

`run_sjsdm_selected_candidate_folds`, [1](#)