

run_sjsdm_tuning_fold_candidates.R

August 12, 2026

Contents

run_sjsdm_tuning_fold_candidates	1
Index	3

run_sjsdm_tuning_fold_candidates
<i>Run sjSDM Tuning Candidates for One Fold</i>

Description

Prepares one fold once, then fits, predicts, and scores every regularization candidate while preserving failures as compact status rows.

Usage

```
run_sjsdm_tuning_fold_candidates(  
  data_candidates = NULL,  
  list_fold_context = NULL,  
  prepare_fold_function = NULL,  
  fit_function = NULL,  
  predict_function = NULL,  
  score_function = score_sjsdm_tuning_predictions,  
  seed = 900723L,  
  epsilon = 1e-06,  
  retain_prediction_cache = FALSE  
)
```

Arguments

`data_candidates`
Candidate table returned by `[build_sjsdm_regularization_candidates()]`.

`list_fold_context`
Fold metadata returned by `[build_sjsdm_tuning_fold_context()]`.

`prepare_fold_function`, `fit_function`, `predict_function`, `score_function`
Injectable fold preparation, fit, prediction, and scoring functions documented by `[run_sjsdm_tuning_candidates()]`.

<code>seed</code>	Non-negative base integer used to derive candidate fit and score seeds.
<code>epsilon</code>	Probability clipping tolerance passed to <code>[score_sjsdm_tuning_predictions()]</code> .
<code>retain_prediction_cache</code>	Logical. When true, return compact prepared-fold metadata and candidate probability matrices alongside the unchanged tuning table.

Value

Compact candidate-by-fold tuning table with metrics and structured status. When ‘`retain_prediction_cache`’ is true, returns a named list containing that table and a compact fold prediction cache without fitted model objects.

Index

`run_sjsdm_tuning_fold_candidates`, [1](#)