

select_sjsdm_regularization.R

September 23, 2026

Contents

select_sjsdm_regularization	1
Index	3

select_sjsdm_regularization
<i>Select sjSDM Regularization</i>

Description

Selects one regularization candidate from repeat-level tuning summaries.

Usage

```
select_sjsdm_regularization(  
  data_tuning_summary = NULL,  
  selection_metric = "negative_log_likelihood_per_response",  
  allow_empty = FALSE  
)
```

Arguments

data_tuning_summary	Repeat-by-candidate table returned by [summarise_sjsdm_tuning_candidates()].
selection_metric	Metric minimized during selection. Defaults to normalized held-out negative log likelihood, "negative_log_likelihood_per_response".
allow_empty	Logical. When 'TRUE', return the typed-empty unit selection if no candidate completed every repeat so the caller can use tier-pooled regularization.

Details

A candidate is eligible only when every represented repeat is complete. The selection metric is averaged equally across repeats. Lower values are preferred, and exact ties are resolved by lexical candidate identifier.

Value

One-row tibble containing the selected candidate parameters, criterion name and value, repeat count, deterministic candidate rank, and regularization source.

Examples

```
## Not run:  
select_sjsdm_regularization(data_tuning_summary = data_summary)  
  
## End(Not run)
```

Index

`select_sjsdm_regularization`, [1](#)