

summarise_sjsdm_model_provenance.R

August 16, 2026

Contents

summarise_sjsdm_model_provenance	1
Index	3

summarise_sjsdm_model_provenance
<i>Summarise sjSDM Model Provenance</i>

Description

Combines cross-validation feasibility, selected regularization, and selected-fold diagnostics into one model-level provenance row.

Usage

```
summarise_sjsdm_model_provenance(  
  data_feasibility = NULL,  
  data_regularization = NULL,  
  data_fold_diagnostics = NULL,  
  fit_device = NULL  
)
```

Arguments

- | | |
|-----------------------|--|
| data_feasibility | One-row cross-validation feasibility table. |
| data_regularization | One-row selected regularization table returned by [resolve_sjsdm_regularization_for_fit()]. |
| data_fold_diagnostics | Selected-candidate fold diagnostics. An empty data frame is allowed when held-out evaluation is unavailable. |
| fit_device | Scalar fitting-device identifier. Must be "cpu" or "gpu". |

Value

One-row tibble containing model context, feasibility, data counts, fold-fit counts, effective MEV minimum/maximum/status and retained-taxon counts, and regularization source. It also records the fitting device and versioned fold-local evaluation contract. The legacy scalar ‘n_effective_mev’ is retained when the count is constant across folds and is missing when counts vary.

Examples

```
## Not run:
summarise_sjsdm_model_provenance(
  data_feasibility = data_cross_validation_feasibility,
  data_regularization = model_regularization_for_fit,
  data_fold_diagnostics = data_sjsdm_out_of_fold_diagnostics,
  fit_device = "gpu"
)

## End(Not run)
```

Index

`summarise_sjsdm_model_provenance`, [1](#)